

Collaborative Filtering

ECON 4810

General Procedure to Model Fitting

1. Hyperparameters
2. Initialize parameters
3. Main loop
 - a. Fit parameters
 - b. Update loss

SVD with Alternating Least Squares

1. Hyperparameters

k = No. of latent factors

2. Initialize parameters

$P = \text{random}(n, k)$
 $Q^T = \text{random}(k, i)$
 $\hat{X} = PQ^T$

These parts
change when the
model change

$\text{loss_prev} = \text{mean}[(X - \hat{X})^2]$

3. Main loop

while loss_delta is not zero:

a. Fit parameters with least squares

$Q^T = (P^T P)^{-1} P^T X$
 $P = X Q (Q^T Q)^{-1}$
 $\hat{X} = P Q^T$

This part changes
when the *method*
of fitting the
model change

b. Update loss

$\text{loss} = \text{mean}[(X - \hat{X})^2]$

$\text{loss_delta} = \text{loss} - \text{loss_prev}$

$\text{loss_prev} = \text{loss}$

SVD with Alternating Least Squares

- SVD with ALS is very fast, typically converge within a dozen epochs.
- The use of least squares means the process is easily parallelized, so scaling up is not a problem
- ALS might not be feasible if you want $X \neq PQ^T$.

Gradient Descent

- Gradient descent is a very general optimization algorithm, usable with pretty much all models.
- It uses the first derivative of the loss function with respect to a parameter to adjust the parameter.

Gradient Descent

- Gradient descent uses the first derivative of the loss function with respect to a parameter to adjust the parameter.
- E.g. suppose our model is

$$\hat{y} = \alpha + x$$

This is a (simplified) regression task, which we can solve by minimizing the squared error:

$$c = (y - \hat{y})^2$$

- For illustration purpose, let us assume the data is generated by $y = 5 + x$, so that if $x = 1, y = 6$.

Gradient Descent


$$\text{Error } \epsilon = y - \hat{y}$$

$$\text{Loss} = \epsilon^2$$

First derivative of loss function:

$$\begin{aligned} \frac{d}{d\alpha} \epsilon^2 \\ &= 2\epsilon \frac{d}{d\alpha} [y - (\alpha + x)] \\ &= -2\epsilon \end{aligned}$$

This is the marginal effect, or **gradient**, of α on loss.



Gradient Descent

We have an initial guess of what α is (usually just a random number.)

This gives us an initial prediction \hat{y} and corresponding error and loss.


e.g.

$$\hat{\alpha}_0 = 2$$

$$\hat{y}_0 = 2 + 1 = 3$$

$$\epsilon_0 = y - \hat{y} = 6 - 3 = 3$$

$$\frac{d}{d\alpha} \epsilon^2 = -2\epsilon = -6$$




Gradient Descent

$$\epsilon_0 = 3$$

$$\frac{d}{d\alpha} \epsilon^2 = -6 > 0$$

In our example, ϵ is positive, so the gradient is negative.

This means loss is decreasing in alpha.




Gradient Descent

This makes sense,
because if

$$y > \hat{y} = \alpha + x$$

increasing α will bring
 \hat{y} closer to y .




Gradient Descent

Conversely, if $\epsilon < 0$,
then

$$y < \hat{y} = \alpha + x$$

decreasing α will bring
 \hat{y} closer to y .



Gradient Descent


The gradient tells us the direction we need to adjust our parameter.

The amount we need to adjust is, to a first-order approximation, inversely proportional to the gradient.

E.g.

$$\frac{d}{d\alpha} \epsilon^2 = -6$$

α needs to be bigger.




Gradient Descent

If we adjust the parameter by the exact amount of the gradient, we might overshoot:

$$\begin{aligned}\hat{\alpha}_1 &= \hat{\alpha}_0 - \frac{d}{d\alpha} \epsilon^2 \\ &= 2 - (-6) \\ &= 8 > 5 = \text{true } \alpha\end{aligned}$$

Overshooting could result in our parameters bouncing around the true value, never to converge.



Gradient Descent


To prevent overshooting, we moderate the gradient by a *learning rate* γ :

$$\hat{\alpha}_1 = \hat{\alpha}_0 - \gamma \frac{d}{d\alpha} \epsilon^2$$

E.g. if $\gamma = 0.1$:

$$\begin{aligned} \hat{\alpha}_1 &= \hat{\alpha}_0 - \gamma \frac{d}{d\alpha} \epsilon^2 \\ &= 2 - 0.1(-6) \\ &= 2.6 < 5 = \text{true } \alpha \end{aligned}$$

Now we are not overshooting, but it will take us more iterations to get to the true α value.



Gradient Descent

- The gradient tells us the direction we need to adjust our parameter. In the above example, the amount we need to adjust is, to a first-order approximation, proportional to $-dc/d\alpha$.
- Moderating the adjustment by a learning rate γ , we have the following update rule:

$$\alpha_t = \alpha_{t-1} - \gamma \left. \frac{dc}{d\alpha} \right|_{\alpha_{t-1}}$$


Or more typical in computer science:

$$\alpha \leftarrow \alpha - \gamma \frac{dc}{d\alpha}$$

- With more than one sample, we take the average.

SVD and Gradient Descent

- In SVD, we have $\hat{X} = PQ^T$
- User u 's preference for item i is $\hat{x}_{ui} = \vec{p}_u \vec{q}_i^T$
- Loss function is $L = \sum \epsilon_{ui}^2 = \sum (\vec{x}_{ui} - \hat{x}_{ui})^2$
- Take derivative of loss function w.r.t. to \vec{p}_u and \vec{q}_i^T (note that they are both vectors):



$$\frac{\partial L}{\partial p_{u,1}} = 2\epsilon_{ui}(-q_{i,1}), \frac{\partial L}{\partial p_{u,2}} = 2\epsilon_{ui}(-q_{i,2}), \dots$$

$$\text{So we have } \frac{\partial L}{\partial \vec{p}_u} = 2\epsilon_{ui}(-\vec{q}_i^T) = -2\epsilon_{ui}\vec{q}_i$$

$$\text{Similarly, } \frac{\partial L}{\partial \vec{q}_i} = 2\epsilon_{ui}(-\vec{p}_u) = -2\epsilon_{ui}\vec{p}_u$$

- Update rule is $\vec{p}_u = \vec{p}_u - \gamma(-2\epsilon_{ui}\vec{q}_i)$, $\vec{q}_i = \vec{q}_i - \gamma(-2\epsilon_{ui}\vec{p}_u)$
- Combine every u and i , we have $P = P + \gamma 2EQ$ and $Q^T = Q^T + \gamma 2P^TE$

SVD with Gradient Descent

1. Hyperparameter

k = No. of latent factors

2. Random initial values

$P = \text{random}(n, k)$
 $Q^T = \text{random}(k, i)$
 $\hat{X} = PQ^T$

$\text{loss_prev} = \text{mean}[(X - \hat{X})^2]$

3. Main loop

while loss_delta is not zero:

a. Fit parameters with
gradient descent

$E = X - \hat{X}$
 $Q^T = Q^T + \gamma P^T E$
 $P = P + \gamma EQ$
 $\hat{X} = PQ^T$

b. Update loss

$\text{loss} = \text{mean}[(X - \hat{X})^2]$
 $\text{loss_delta} = \text{loss} - \text{loss_prev}$
 $\text{loss_prev} = \text{loss}$

Simon Funk's SVD

- User u 's preference for item i is $\hat{x}_{ui} = \mu + b_u + b_i + \vec{p}_u \vec{q}_i^T$, where μ the constant term, b_u is the user fixed effect and b_i the item fixed effect.
- Compute gradients:

$$\frac{\partial L}{\partial \vec{p}_u} = 2\epsilon_{ui}(-\vec{q}_i^T) = -2\epsilon_{ui}\vec{q}_i$$

$$\frac{\partial L}{\partial \vec{q}_i} = 2\epsilon_{ui}(-\vec{p}_u) = -2\epsilon_{ui}\vec{p}_u$$

$$\frac{\partial L}{\partial \mu} = \frac{\partial L}{\partial b_u} = \frac{\partial L}{\partial b_i} = 2\epsilon_{ui}(-1) = -2\epsilon_{ui}$$

- Update rule is

$$\vec{p}_u = \vec{p}_u - \gamma(-2\epsilon_{ui}\vec{q}_i), \vec{q}_i = \vec{q}_i - \gamma(-2\epsilon_{ui}\vec{p}_u)$$

$$\mu = \mu - \gamma(-2\epsilon_{ui}), b_u = b_u - \gamma(-2\epsilon_{ui}), b_i = b_i - \gamma(-2\epsilon_{ui})$$

Stochastic Gradient Descent

- Learning is quite slow if we only update model weights after we go through all data.
- We could instead update every time after we gone through a given number of samples. This is **Stochastic Gradient Descent (SGD)**.
- Because we are not using all data, we could be updating towards the wrong direction sometimes, but on average the updates will be correct. Hence, *stochastic*.

Stochastic Gradient Descent

- The stochastic nature is actually a good property because it helps the model escape from local optima.